

Lecture 4: Containers II

LIVE
TUTORIAL
BY YOU AND SHIVAS!

AI-5

Productionizing AI (MLOps)

LOADS
OF
FUN!

Pavlos Protopapas, Shivas Jayaram

WITH FUNKY SLIDES BY
RAHUL DAVE

RECAP / MICROSERVICES /
DOCKER BUILDS / LIVE
PIPELINE!

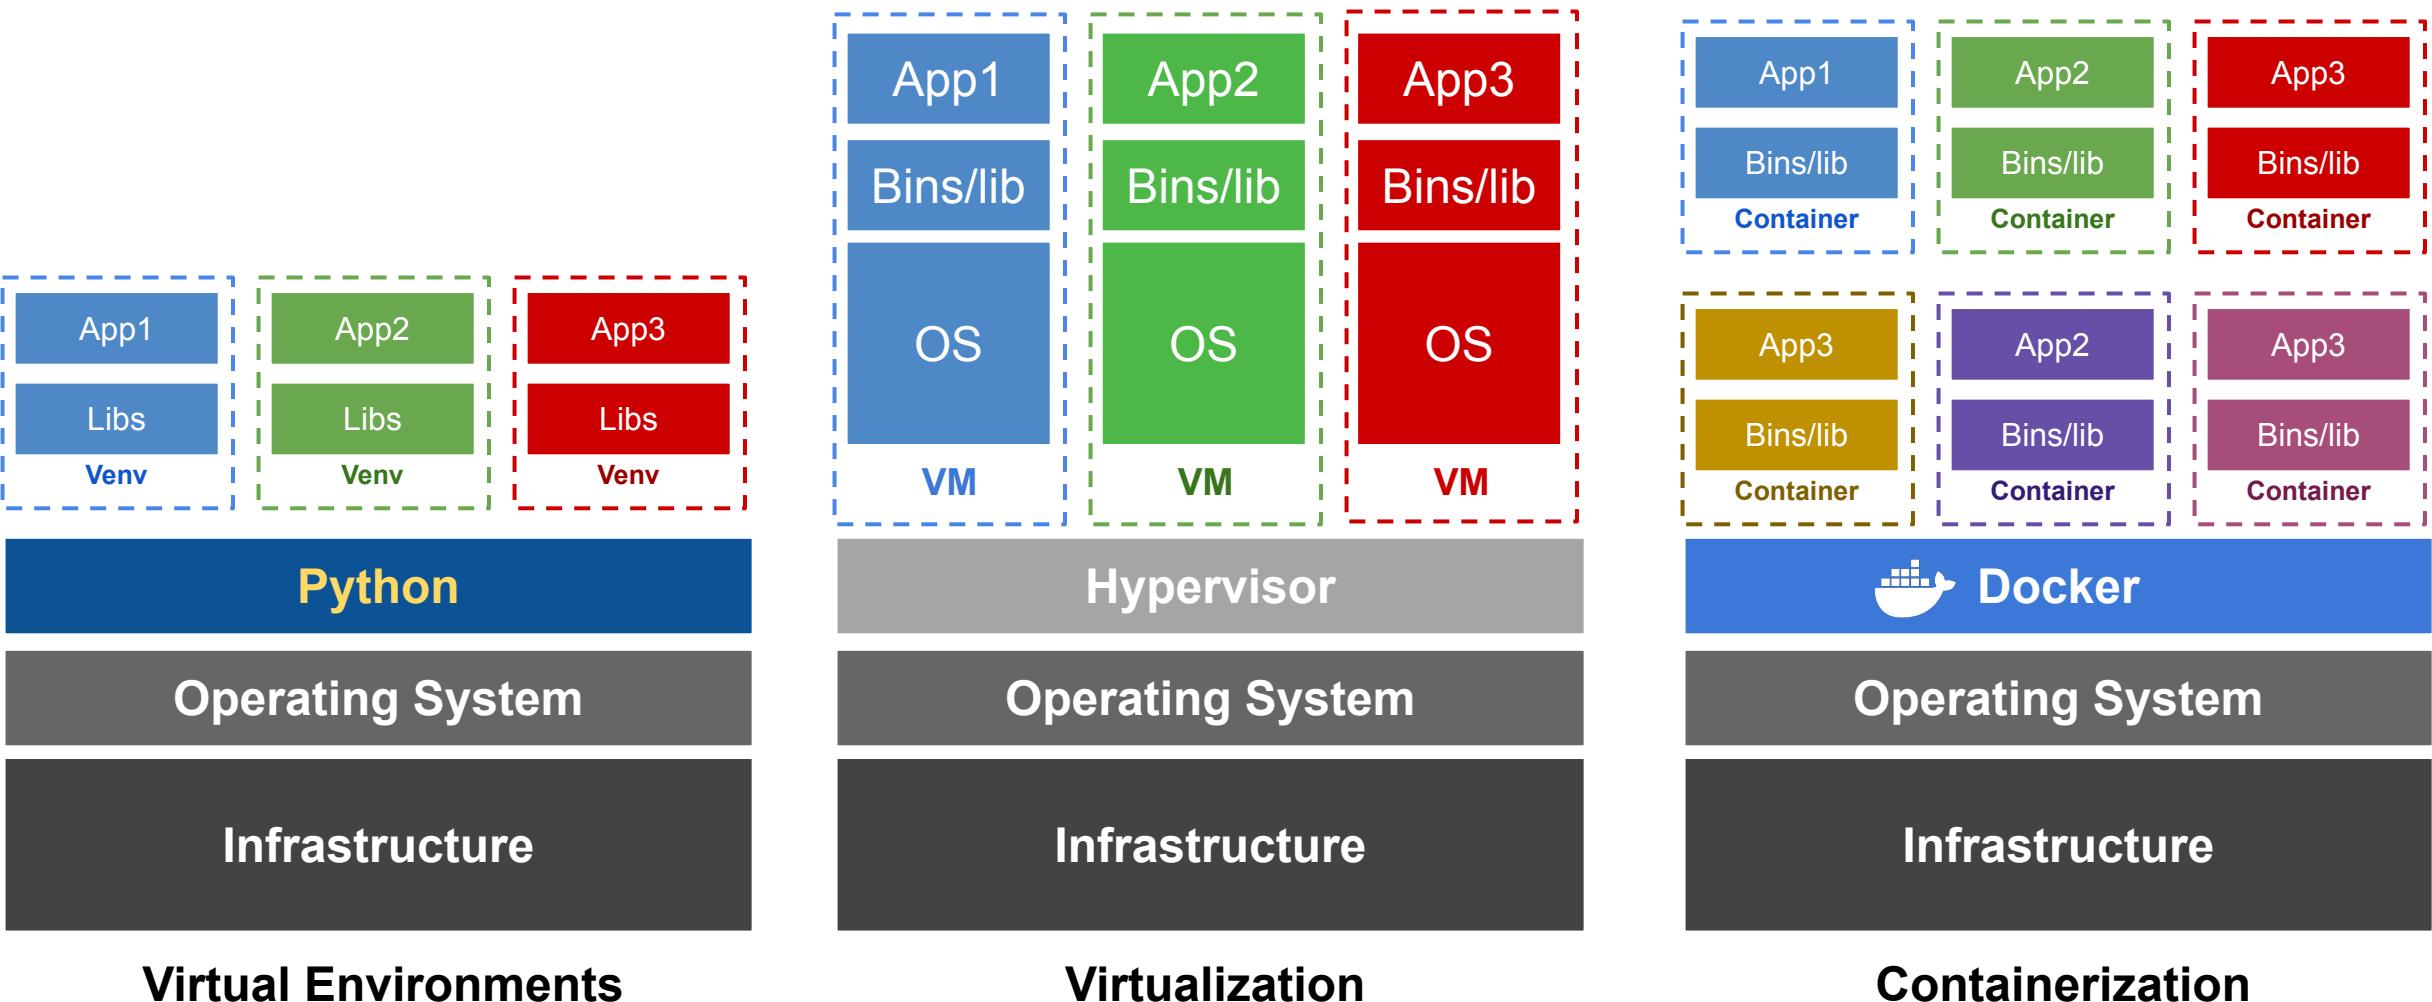
Outline

1. Recap: Review of Previous Material
2. Containers in Architecture: Microservices vs. Monolithic
3. Implementing Containers as Microservices

Outline

- 1. Recap: Review of Previous Material**
2. Containers in Architecture: Microservices vs. Monolithic
3. Implementing Containers as Microservices

Recap: Environments vs Virtualization vs Containerization

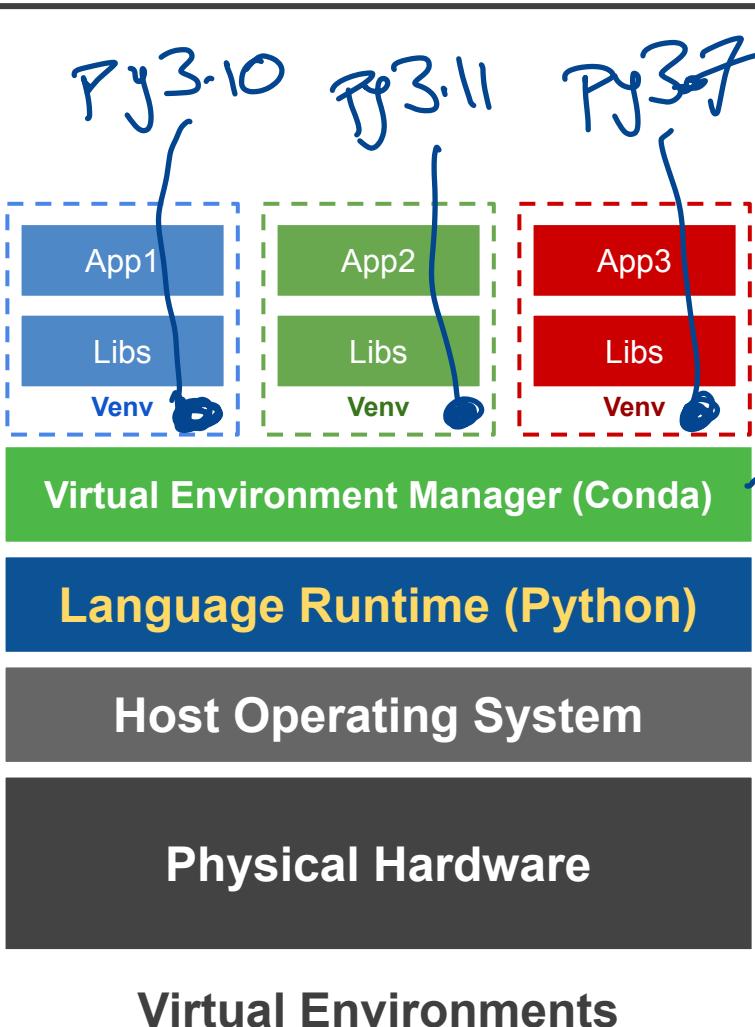


Is This Really the right question?

Real world scenarios will often combine.

- VM's are great to isolate tenants from each other
- But containers are best to isolate "jobs" from each other
- Virtual Environments (venvs) are useful in either case. Indeed, they make building containers EASIER!

Environments



Also venv, Pipenv, conda-lock, poetry, Pixi,

- **Dependency Isolation:** Virtual environments modify the PATH and other environment variables so that the dependencies are loaded from the environment's directory, rather than system-wide directories
- **No Kernel Isolation:** Unlike VMs and containers, virtual environments don't provide any kernel level isolation.
- **Resource Utilization:** Since virtual environments don't have any additional OS or kernel, they are the most efficient in terms of resource utilization among the three.
- **Filesystem Boundaries:** Virtual environments usually don't provide isolation at the filesystem level; files written in one environment are accessible from others.

CONDA VS PIPENV : THE LOCKFILE ISSUE

CONDA

- NON PYTHON SOFTWARE
- CUDA / HARDWARE

PIPEENV / POETRY / ETC.

- ONLY PYTHON SOFTWARE
- CUDA IS HARDER

PEOPLE OFTEN COMBINE

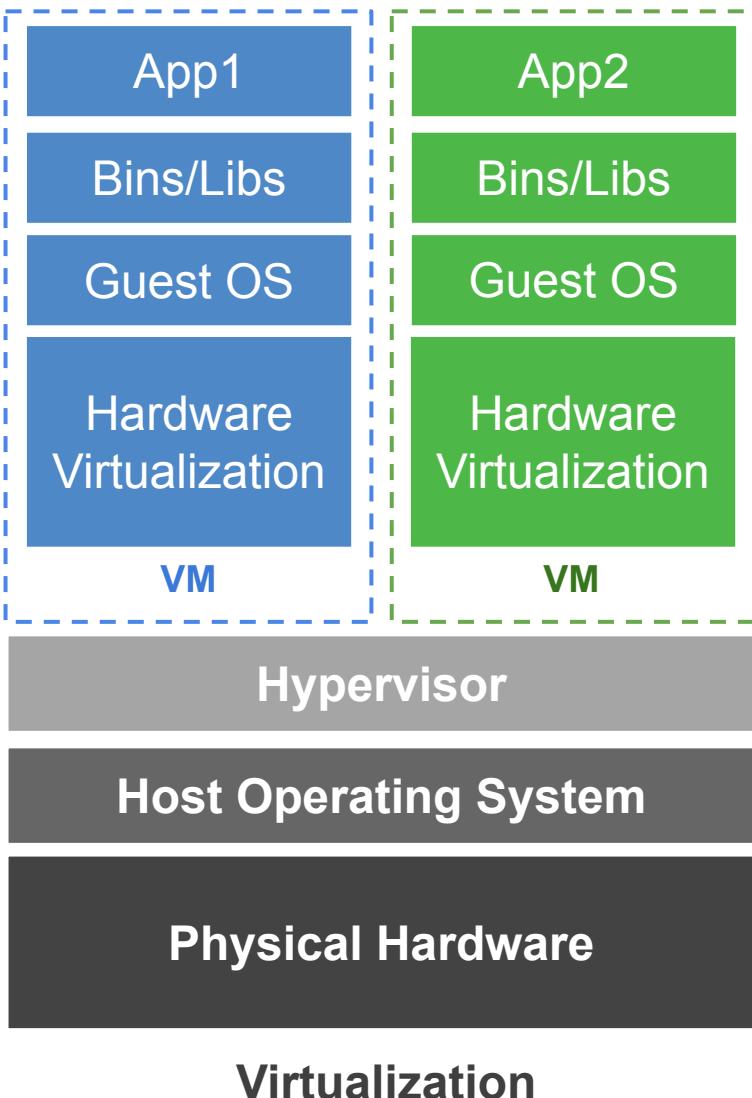
conda for cuda / python exec | rust | rkt++ | jcuva [conda-lock]
pip / poetry | pip-tools [latter 2 provide lock]

PROBLEM: lack of combined lockfiles.

SOLUTION: no general one yet, pip combines
lockfiles

EASIER TO STAY WITH ONE, BUT U DONT ALWAYS HAVE THE CHOICE

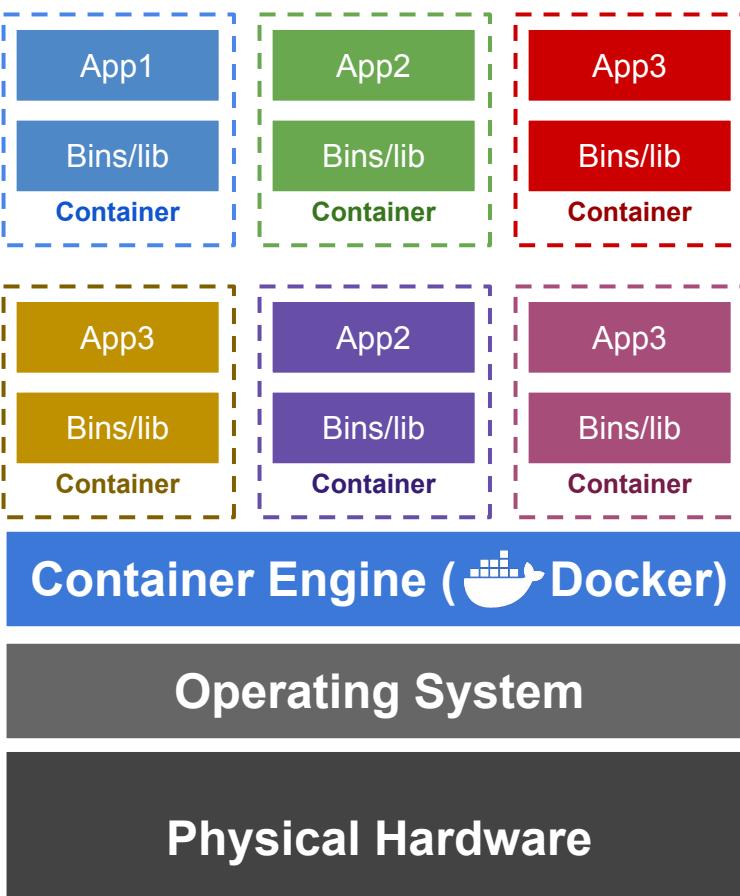
Virtualization (Virtual Machines)



- **CPU Virtualization:** VMs usually have a set number of virtual CPU cores allocated by the hypervisor. These virtual CPUs map to physical CPU cores, but the hypervisor adds a layer of management and overhead, which can lead to inefficiencies.
- **Emulated Devices:** VMs have emulated hardware devices, meaning the VM sees virtual CPUs, virtual network adapters, and virtual disks that the hypervisor translates to real hardware resources.
- **Full OS:** Each VM runs its full guest OS. This means that each VM has its own separate kernel space and user space, making resource management fully independent but less efficient.
- **Resource Allocation:** RAM and CPU are often (not always) allocated in blocks, and disk space is generally pre-allocated, making VMs less flexible in terms of resource utilization.

Containerization

REPRODUCIBILITY TO THE OS LEVEL



- **Namespaces**: Containers use kernel features like namespaces to provide isolation of processes and resources. This allows each container to operate as if it is the only application running on the system. Example namespaces include:
 - PID Namespace: Isolates the process ID number space. In other words, processes in different PID namespaces can have the same PID.
 - Mount Names: Isolates the file system tree so that each namespace can have its own file system layout.
- **Control Groups (cgroups)**: Complementary to namespaces, cgroups limit resource usage, like CPU, memory, and IO, allowing for better resource utilization compared to VMs.
- **Process Virtualization**: Namespaces and cgroups together enable process virtualization by allowing processes to run in isolated environments with controlled access to system resources.
- **Shared Kernel**: Containers share the host's OS kernel but have their own filesystem, libraries, and bins, making them lightweight yet isolated.
- **Direct Access**: Containers can access host resources more directly, avoiding much of the overhead introduced by hypervisors in VMs.

EXAMPLE DOCKERFILE

What is
bad about
this?

Layers

```
# Use the official Debian-hosted Python image
FROM python:3.9-slim-buster

# Tell pipenv where the shell is.
# This allows us to use "pipenv shell" as a container entry point.
ENV PYENV_SHELL=/bin/bash

# Ensure we have an up to date baseline, install dependencies
RUN set -ex; \
    apt-get update && \
    apt-get upgrade -y && \
    apt-get install -y --no-install-recommends build-essential git && \
    pip install --no-cache-dir --upgrade pip && \
    pip install pipenv

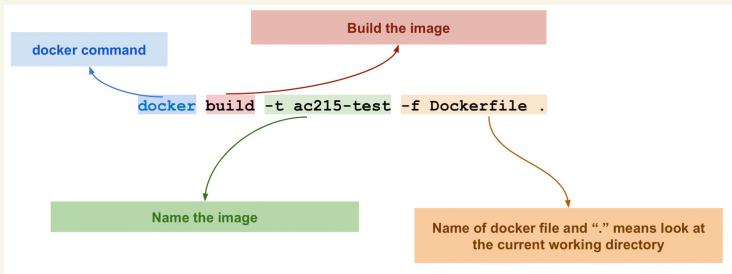
# Add Pipfile, Pipfile.lock + python code
ADD . /
RUN pipenv sync

# Entry point
ENTRYPOINT ["/bin/bash"]

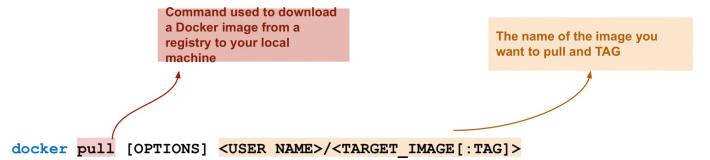
# Get into the pipenv shell
CMD ["-c", "pipenv shell"]
```

SOME DOCKER COMMANDS

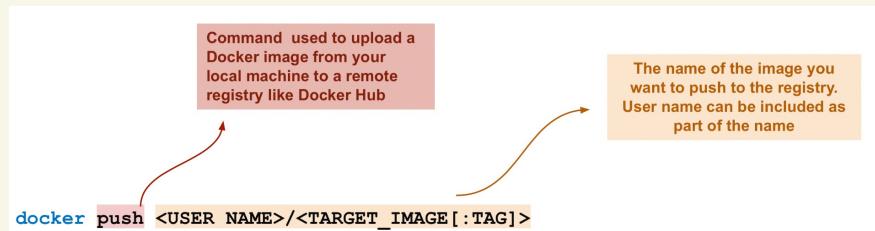
BUILD



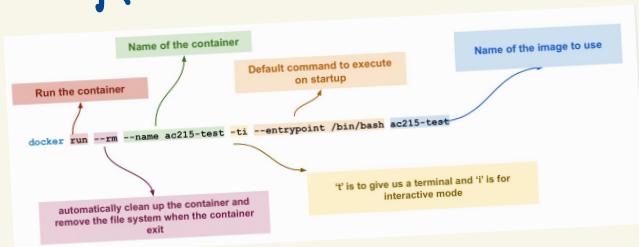
PULL



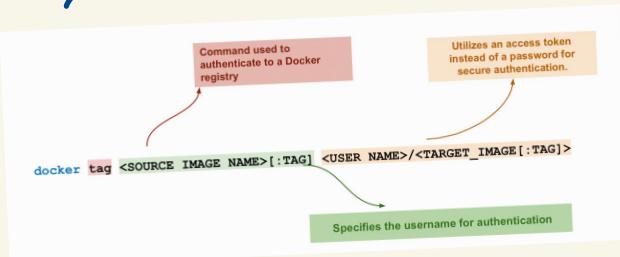
PUSH



RUN



TAG



DOCKERFILE

VS IMAGE

VS

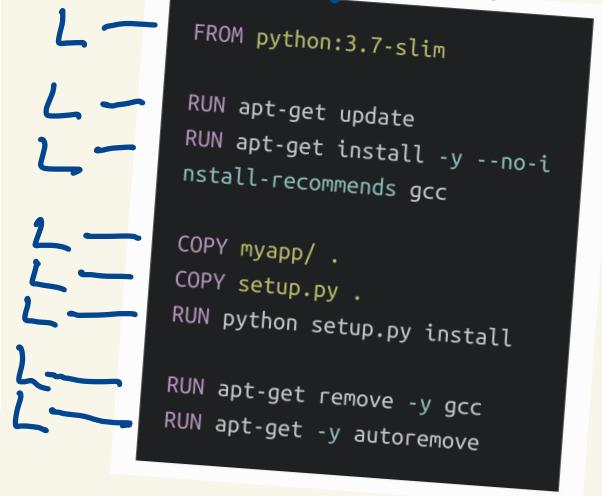
CONTAINER

```
FROM python:3.7-slim

RUN apt-get update
RUN apt-get install -y --no-
install-recommends gcc

COPY myapp/ .
COPY setup.py .

RUN python setup.py install
```

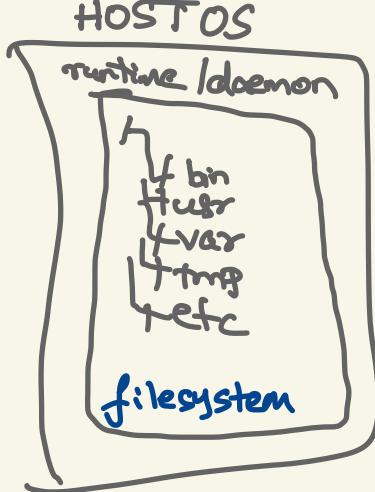


Just a file!

The image built
243 MB

Layers are cached on host

This Image
245 MB



A not so large container

PROBLEM DOCKER IMAGES ARE APPEND-ONLY

- The image has layers with gcc in them to build our project
- It also has a remove layer
This does NOT REDUCE image size
- The container has the files removed
BUT

YOU HAVE A LARGE IMAGE AND

YOU HAVE TO DOWNLOAD COMPILER
TO BUILD.

SOLUTION 1 : SINGLE LAYER

```
FROM python:3.7-slim

COPY myapp/ .
COPY setup.py .

RUN apt-get update && \
    apt-get install -y --no-i
    nstall-recommends gcc && \
        python setup.py install
&& \
    apt-get remove -y gcc &&
    apt-get -y autoremove
```

161 MB IMAGE

Put everything into one layer so that complex files are gone

PROBLEM

- we destroyed the build cache
 - builds are slower
- (see middle dockerfile 2 slides back)

or call "docker build -squash"

Running commands can take up a lot of disk space. For example, when **installing** and **building** packages, we **download** and **produce** many files. This can make our image size **very large**.

Question: Does our app need all of these files?

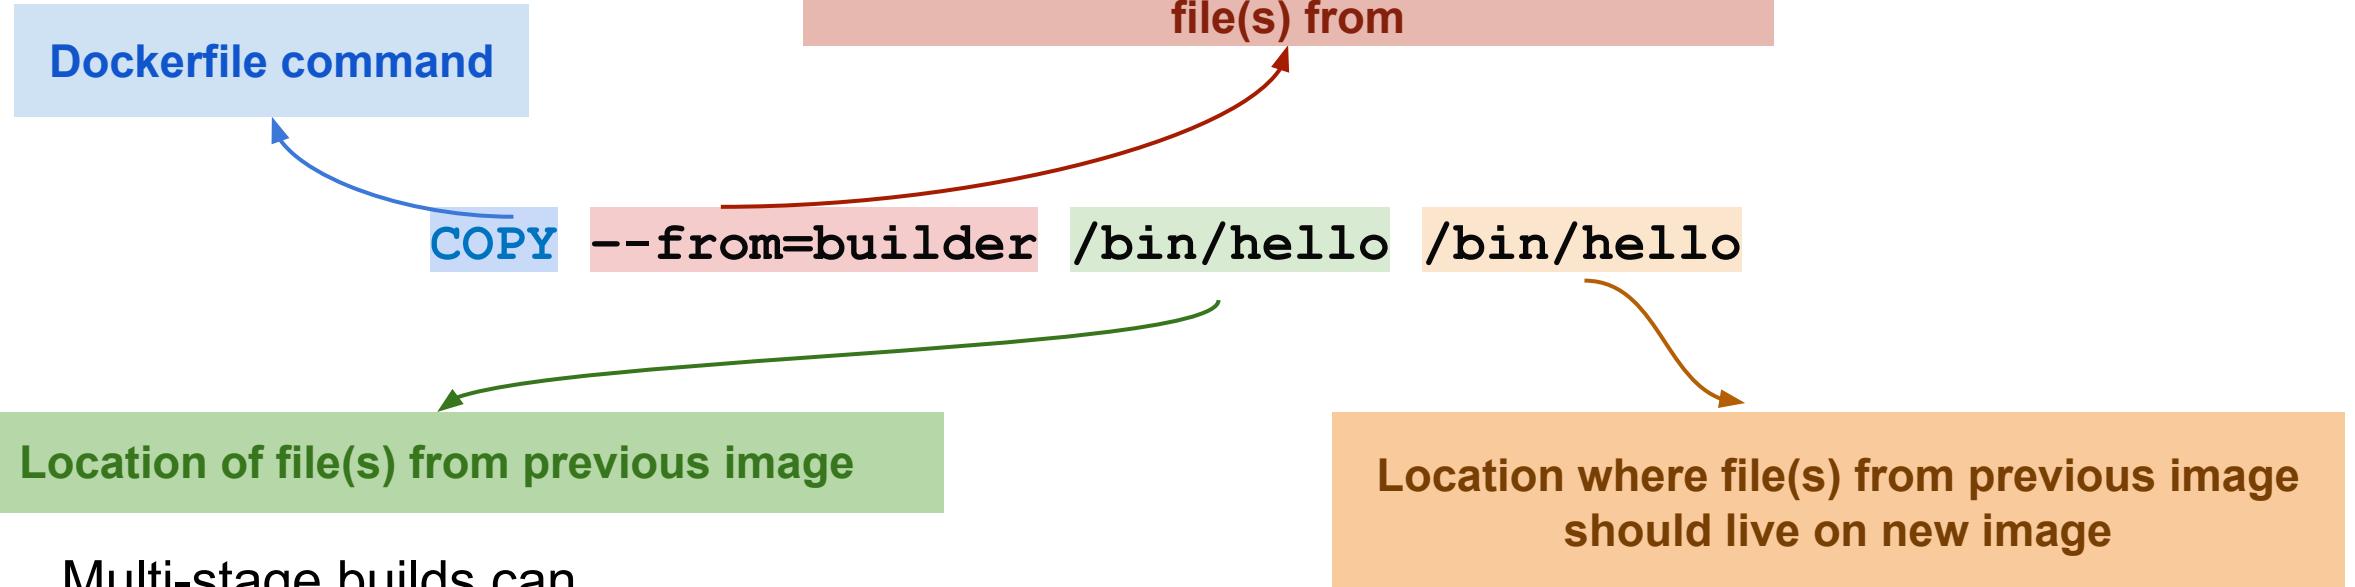
Answer: Usually no! We just need the executable and its runtime dependencies.

Multi-stage builds allow us to bring **only what we need** into the final Docker image.

Pro Tips 1: multi-stage builds

Multi-stage builds in Docker allow the use of multiple **FROM** statements in a single **Dockerfile**.

Each stage can bring in files from the previous stage using the **COPY** instruction.



Multi-stage builds can

- decrease container size
- improve security (e.g., avoid sharing private keys from GitHub)
- leverage different base images for each stage while preserving only the final image.

Pro Tips 1: multi-stage builds: example

```
# First Stage: Building the Python application
FROM python:3.8 AS build-env
WORKDIR /app
COPY BLAH BLAH
RUN BLAH BLAH

# Second Stage: Copy the dependencies and run the application
FROM python:3.8-slim
COPY --from=build-env /usr/local/lib/python3.8/site-packages
/usr/local/lib/python3.8/site-packages
WORKDIR /app
COPY BLAH BLAH
CMD BLAH BLAH
```

```
# This is the first image:  
FROM ubuntu:18.04 AS compile-image  
RUN apt-get update  
RUN apt-get install -y --no-install-recommends gcc build-essential  
  
WORKDIR /root  
COPY hello.c .  
RUN gcc -o helloworld hello.c  
  
# This is the second and final image; it copies the compiled  
# binary over but starts from the base ubuntu:18.04 image.  
FROM ubuntu:18.04 AS runtime-image  
  
COPY --from=compile-image /root/helloworld .  
CMD ["./helloworld"]
```

MULTI STAGE BUILD EXAMPLE

- this is great: small size and no compiler
- **problem 1**: standard python has a problem in that it is splatted all over the filesystem
- One way "pip install --user" and that puts everything into ".local"
- BETTER VENVS

88.9 MB

VENVS : all of Python in one spot

F
I
R
S
T

S
T
A
G
E

```
FROM python:3.7-slim AS compile-image
RUN apt-get update
RUN apt-get install -y --no-install-recommends build-essential gcc

RUN python -m venv /opt/venv
# Make sure we use the virtualenv:
ENV PATH="/opt/venv/bin:$PATH"

COPY requirements.txt .
RUN pip install -r requirements.txt

COPY setup.py .
COPY myapp/
RUN pip install .
```

```
FROM python:3.7-slim AS build-image
COPY --from=compile-image /opt/venv /opt/venv

# Make sure we use the virtualenv:
ENV PATH="/opt/venv/bin:$PATH"
CMD ['myapp']
```

← SECOND STAGE

Virtual Environments
and
Containers
come together
to make multi-
-stage builds
nice.

← or install
from pipenv
lockfile

DOES THIS Fix ALL PROBLEMS? MOSTLY

BUT

①.

No. Multi-Stage Builds can be slow.
Why? Note that gcc layer in first build may not be in cache if you pull from registry. Solution: push/pull all stages.

②

CAVEAT : if you are using a venv system like poetry or anything based on "pyproject.toml", be careful. Adding a new task into the file can invalidate the cache.

STARTING WITH (POSSIBLY BUILT) CONTAINER

The -e treats unset variables as errors

```
#!/bin/bash
set -euo pipefail
# Pull the latest version of
the image, in order to
# populate the build cache:
docker pull itamarst/helloworld
ld || true
# Build the new version:
docker build -t itamarst/helloworld .
# Push the new version:
docker push itamarst/helloworld
ld
```

GOOD SHELL HYGIENE

error if any command in
a pipe fails, and all
errors return a failure code

The || true handles first time

<https://sipb.mit.edu/doc/safe-shell/>

what about for 2
stage pipelines?

STAGE 1

```
#!/bin/bash
set -eo pipefail
# Pull the latest version of
the image, in order to
# populate the build cache. T
his isn't necessary if
# you're using BuildKit.
docker pull itamarst/helloworld:compile-stage || true
docker pull itamarst/helloworld:latest || true

# Build the compile stage:
docker build --target compile
-image \
--build-arg BUILDKIT_INLINE_CACHE=1 \
--cache-from=itamarst/
helloworld:compile-stage \
--tag itamarst/helloworld:compile-stage .
```

STAGE 2

```
# Build the runtime stage, us
ing cached compile stage:
docker build --target runtime
-image \
--build-arg BUILDKIT_INLINE_CACHE=1 \
--cache-from=itamarst/
helloworld:compile-stage \
--cache-from=itamarst/
helloworld:latest \
--tag itamarst/helloworld:latest .

# Push the new versions:
docker push itamarst/helloworld:compile-stage
docker push itamarst/helloworld:latest
```

BUILDING IN 2 STAGES

You can build from your local Dockerfile but this make sure to pull down all images to cache.

```
# Push the new versions:
docker push itamarst/helloworld:compile-stage
docker push itamarst/helloworld:latest
```

TAG STAGES EXPLICITLY
match docker classical behavior.

FROM PYTHONSPEED.

DOCKER vs PYTHON

```
FROM python:3.8-slim-buster
COPY requirements.txt /tmp
RUN pip install -r requirements.txt
COPY . /tmp/myapp
RUN pip install /tmp/myapp
```

① COPY AND INSTALL FIRST
(could use lockfile)

② Then put in your app, so any changes in it don't affect cached layers.

WORKAROUNDS

- see poetry issue 1301 on github [try to use only lockfile]
- no probs on - pupenv.
- on pixi: "pixi --frozen"

PROBLEM

say something here changes.

```
[tool.poetry]
name = "myexample"
version = "0.1.0"
description = ""
authors = ["Itamar Turner-Trauring"]

[tool.poetry.dependencies]
python = "3.6"
Flask = "1.1.2"

# ...
```

```
FROM python:3.8-slim-buster
```

```
WORKDIR /app
```

```
# Install poetry:
RUN pip install poetry
```

```
# Copy in the config files:
COPY pyproject.toml poetry.lock ./
```

```
# Install only dependencies:
RUN poetry install --no-root --no-dev
```

```
# Copy in everything else and install:
COPY .
RUN poetry install --no-dev
```

then
INVALIDATION

GENERAL PROBLEM WITH PYPROJECT.TOML

Expect venv managers to fix.

Pro Tips 2: multi-platform images

When building a more complicated Docker image, there is a small chance the specific platform (OS and CPU architecture) on your machine causes issues when sharing the Docker image with someone on a different machine

Example: Building a complex image on M1 Mac (linux/arm64) and trying to run the image on an older Macbook (linux/amd64)

Error message to look out for

The requested image's platform (linux/arm64) does not match the detected host platform (linux/amd64) and nospecific platform was requested

Solution:

- Use the `--platform` flag within the `FROM` command in your Dockerfile to specify the target OS and CPU architecture for the build output

[multi-platform images](#)

Outline

1. Recap: Review of Previous Material
2. **Containers in Architecture: Microservices vs. Monolithic**
3. Implementing Containers as Microservices

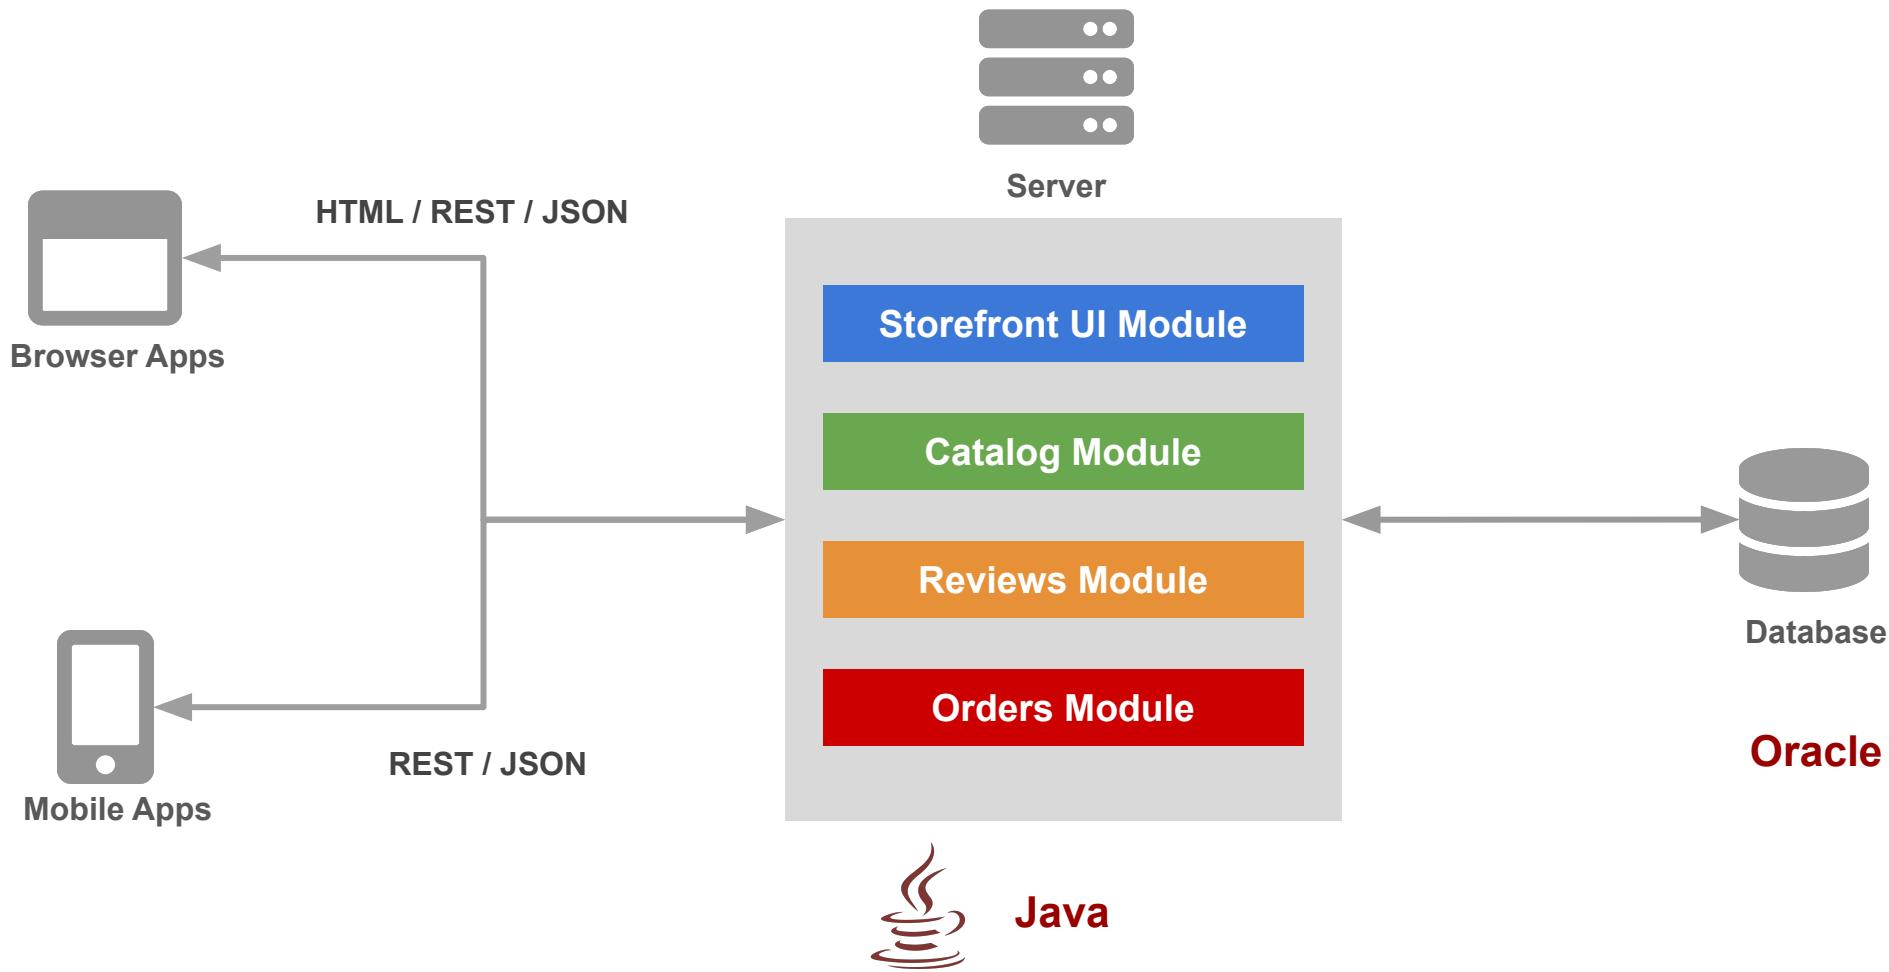
Conceptual Scenario

- Picture building a comprehensive application, such as an online store.

Traditional Approach

- Traditionally you would build this using a Monolithic Architecture

Monolithic Architecture



Monolithic Architecture - Advantages

Simplicity in Development:

Streamlined development process as most tools and IDEs natively support monolithic applications.

Ease of Deployment:

Hassle-free deployment with all components bundled into a single, unified package.

Scalability:

Easier to scale horizontally by replicating the entire application as a whole.

Monolithic Architecture - Disadvantages

Maintenance Challenges:

Complexity increases over time, making it harder to implement changes or find issues.

System Vulnerability:

A failure in a single component can lead to the collapse of the entire system.

Patching Difficulties:

Patching or updating specific modules can be cumbersome due to tightly-coupled components.

Monolithic Architecture - Disadvantages

Technology Lock-in:

Adopting new technologies or updating existing ones can be problematic due to interdependencies.

Slow Startup:

Increased startup time as all components must be initialized simultaneously.

Applications have changed dramatically

A decade ago

Apps were monolithic
Built on a single stack (e.g. .Net or Java)
Long lived
Deployed to a single server

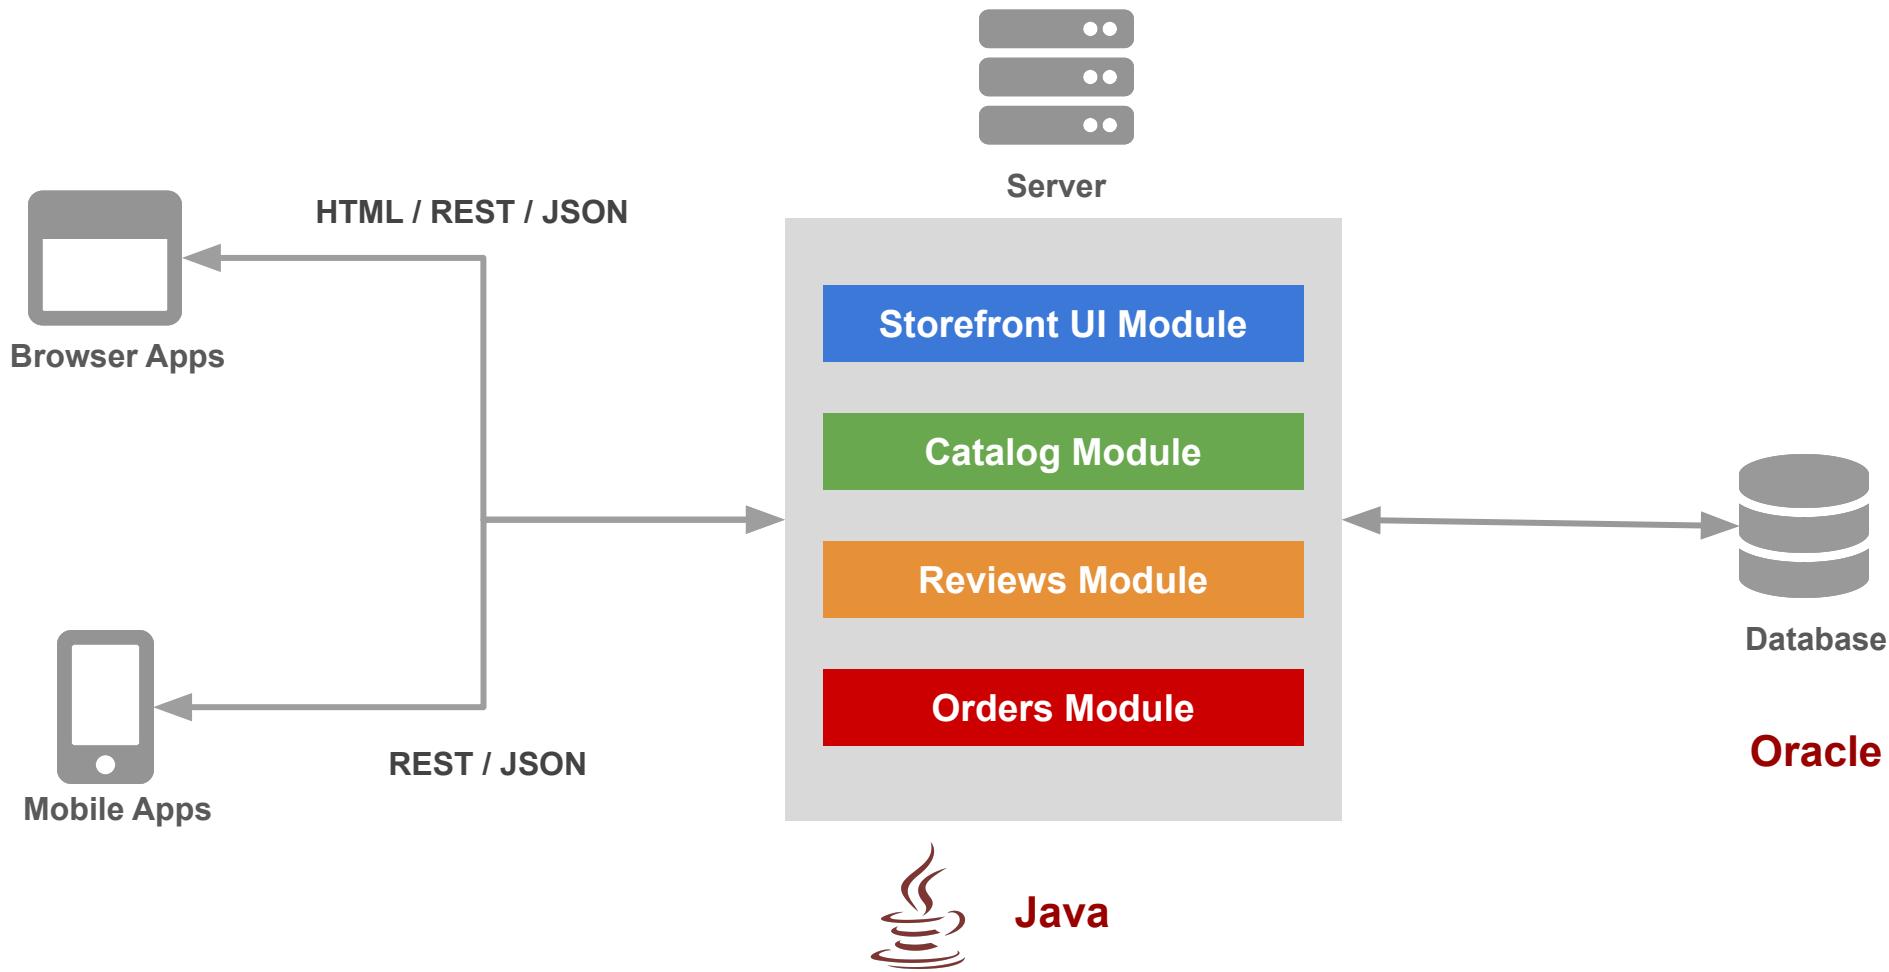
Today

Apps are constantly being developed
Build from loosely coupled components
Newer version are deployed often
Deployed to a multitude of servers

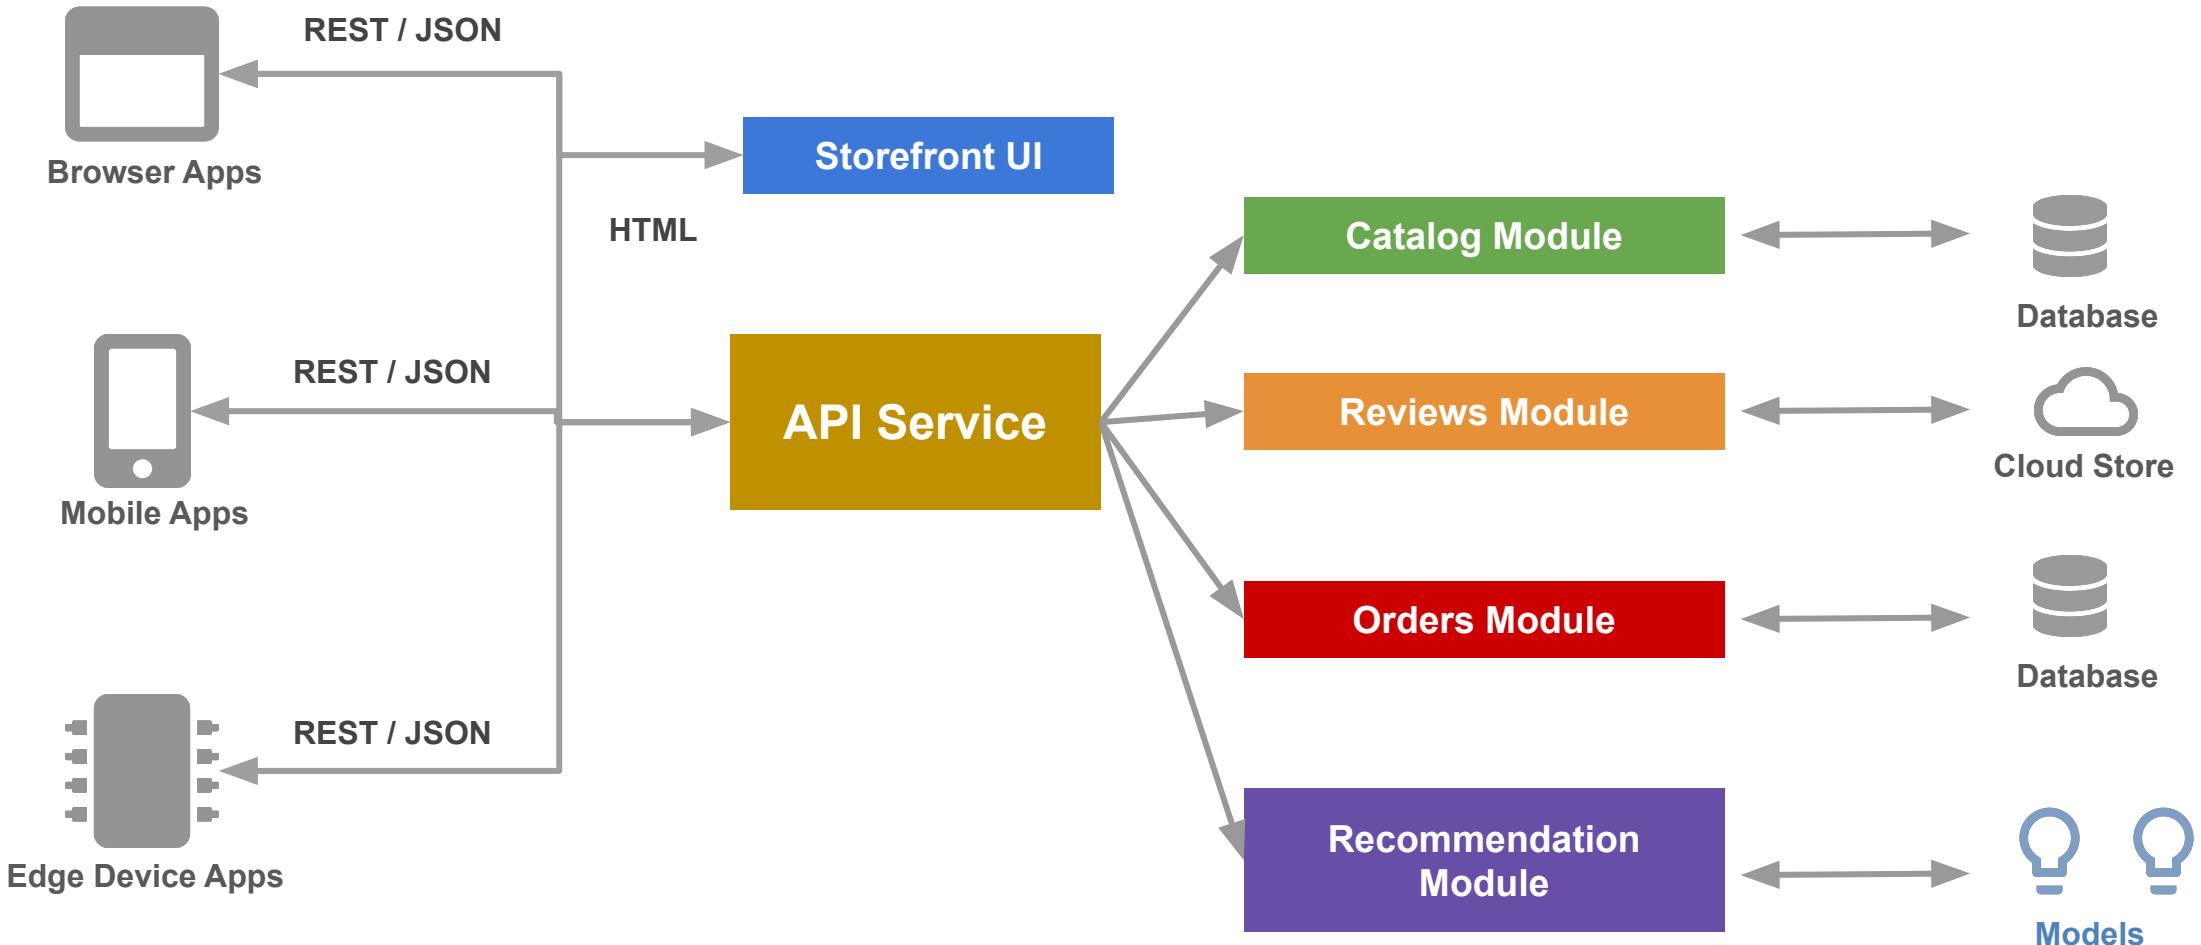
Data Science

Apps are being integrated with various data types/sources and models

Monolithic Architecture



Today: Microservice Architecture



Microservice Architecture - Advantages

Simplified Maintenance:

Modular design makes it easier to manage, update, and debug individual services.

Fault Isolation:

Independent components ensure that failure in one service doesn't bring down the entire application.

Streamlined Patching:

Easier to patch or update specific services without affecting the entire system.

Technological Flexibility:

Adapting to or adopting new technologies becomes seamless due to service independence.

Quick Startup:

Reduced startup time as all components can be initialized in parallel.

Microservice Architecture - Disadvantages

Development Complexity:

Varied technologies across components can complicate the development process.

Deployment Hurdles:

Multiple technologies and dependencies require a complex setup for deployment.

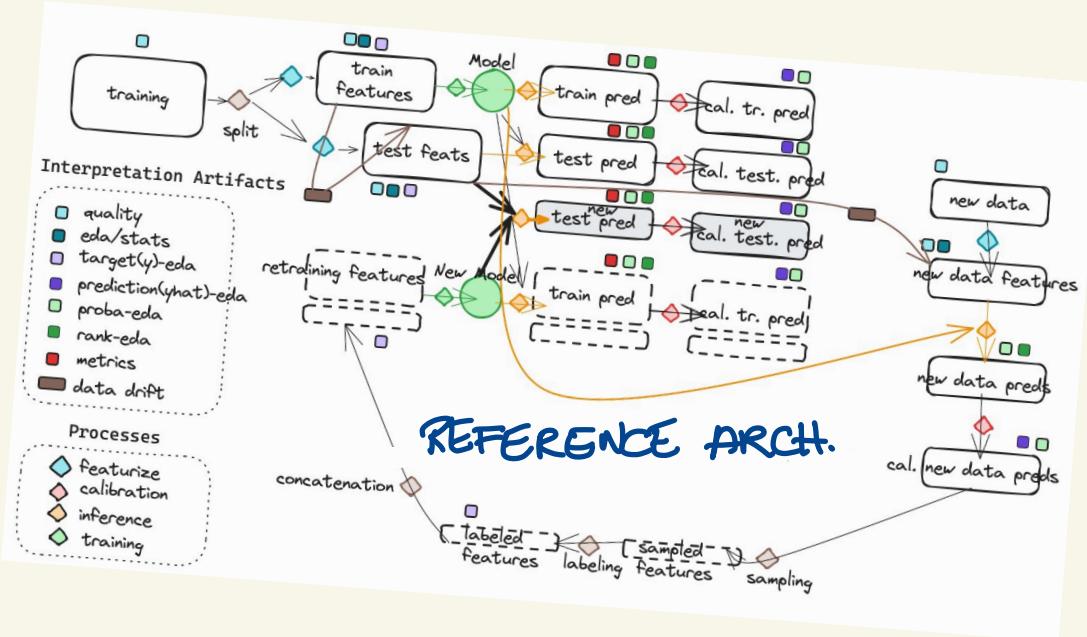
Scaling Concerns:

Scaling the entire application can be intricate due to disparate components.

BUT : Opportunity to Scale Individual Components

Docker + Kubernetes

MICROSERVICES AND DATA SCIENCE



- many moving parts
- data oriented
- processes (eg featureize or inference) repeated in many contexts
- intricate dependencies

- repeated processes can be dockerized as microservices, kept completely general, with filenames, etc passed in as environment. **DOCKER COMPOSE / KUBERNETES**

DATA MEANS STORAGE MEANS MOUNTS

- If you are going to create and trans form data you must store it
- **BIND** mounts are stored anywhere on host system
- **VOLUMES** are stored in part of the host filesystem managed by docker

```
Create a volume named vol-name -  
docker volume create vol-name  
  
View detailed information about the volume, like creation date, mount-point, storage-driver etc. -  
docker volume inspect vol-name  
  
Remove the volume -  
docker volume rm vol-name  
  
Remove all unused local volumes (excluding the ones currently being using by containers) -  
docker volume prune
```

VOLUME

docker run

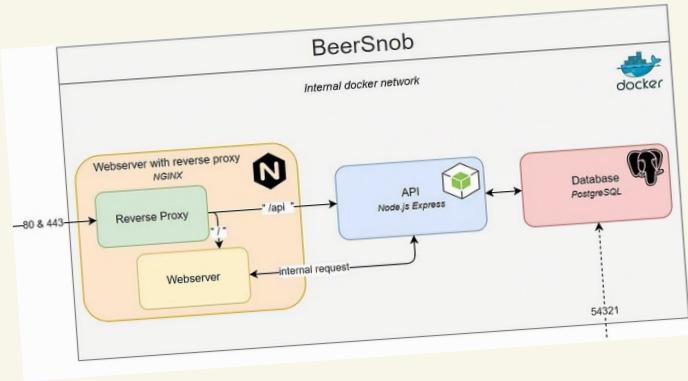
source=\$HOME/docker/volumes/postgr

BIND

--mount type=bind,source="\$pwd",ta

from techmomo.com

DOCKER COMPOSE



```
10 beersnob_api:
11   container_name: beersnob_api
12   hostname: beersnob_api
13   build:
14     context: ./beersnob_api
15     dockerfile: Dockerfile
16   ports:
17     - 54322:5000
18   volumes:
19     - ./beersnob_api/src/:/usr/src/app/
20     - /usr/src/app/node_modules
21   restart: unless-stopped
22   environment:
23     NODE_ENV: ${BEERSNOB_ENVIRONMENT}
24   depends_on:
25     - beersnob_database
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
```

```
57 beersnob_webserver:
58   container_name: beersnob_webserver
59   hostname: beersnob_webserver
60   build:
61     context: ./beersnob_webserver
62     dockerfile: Dockerfile
63   ports:
64     - 80:80
65     - 443:443
66   volumes:
67     - ./beersnob_webserver/src/test:/u
68   restart: always
69   depends_on:
70     - beersnob_database
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
```

```
1 version: '3'
2
3 services:
4
5   beersnob_database:
6     container_name: beersnob_database
7     hostname: beersnob_database
8     image: postgres
9     volumes:
10       - ./volume:/var/lib/postgresql
11     environment:
12       - POSTGRES_DB=beersnobdb, beersnobdb
13       - POSTGRES_USER=mhuls
14       - POSTGRES_PASSWORD=aStrongPassword
15     ports:
16       - 54321:5432
17     restart: unless-stopped
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
39
40
41
42
43
44
45
46
47
48
49
49
50
51
52
53
54
55
56
57
58
59
59
60
61
62
63
64
65
66
67
68
69
69
70
71
72
73
74
75
76
77
78
79
79
80
81
82
83
84
85
86
87
88
89
89
90
91
92
93
94
95
96
97
98
99
```

<https://towardsdatascience.com/docker-compose-for-absolute-beginners-how-does-it-work-and-how-to-use-it-examples-733ca24c5e6c>

Why use Containers?

DEVELOPER EXPERIENCE

- Consider a software development team workflow for developing an App
- Traditionally you would develop/build this independently in various machines (dev, test, qa, prod)

Software Development Workflow (no Docker)

Windows



Node.js
Python



Linux



Node.js
Python



Mac



Node.js
Python



OS Specific **installation** in
every developer machine

Source Control

GitHub

Build Server

Linux

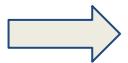


Production / Test Servers

Linux



Linux



Every team member moves
code to source control

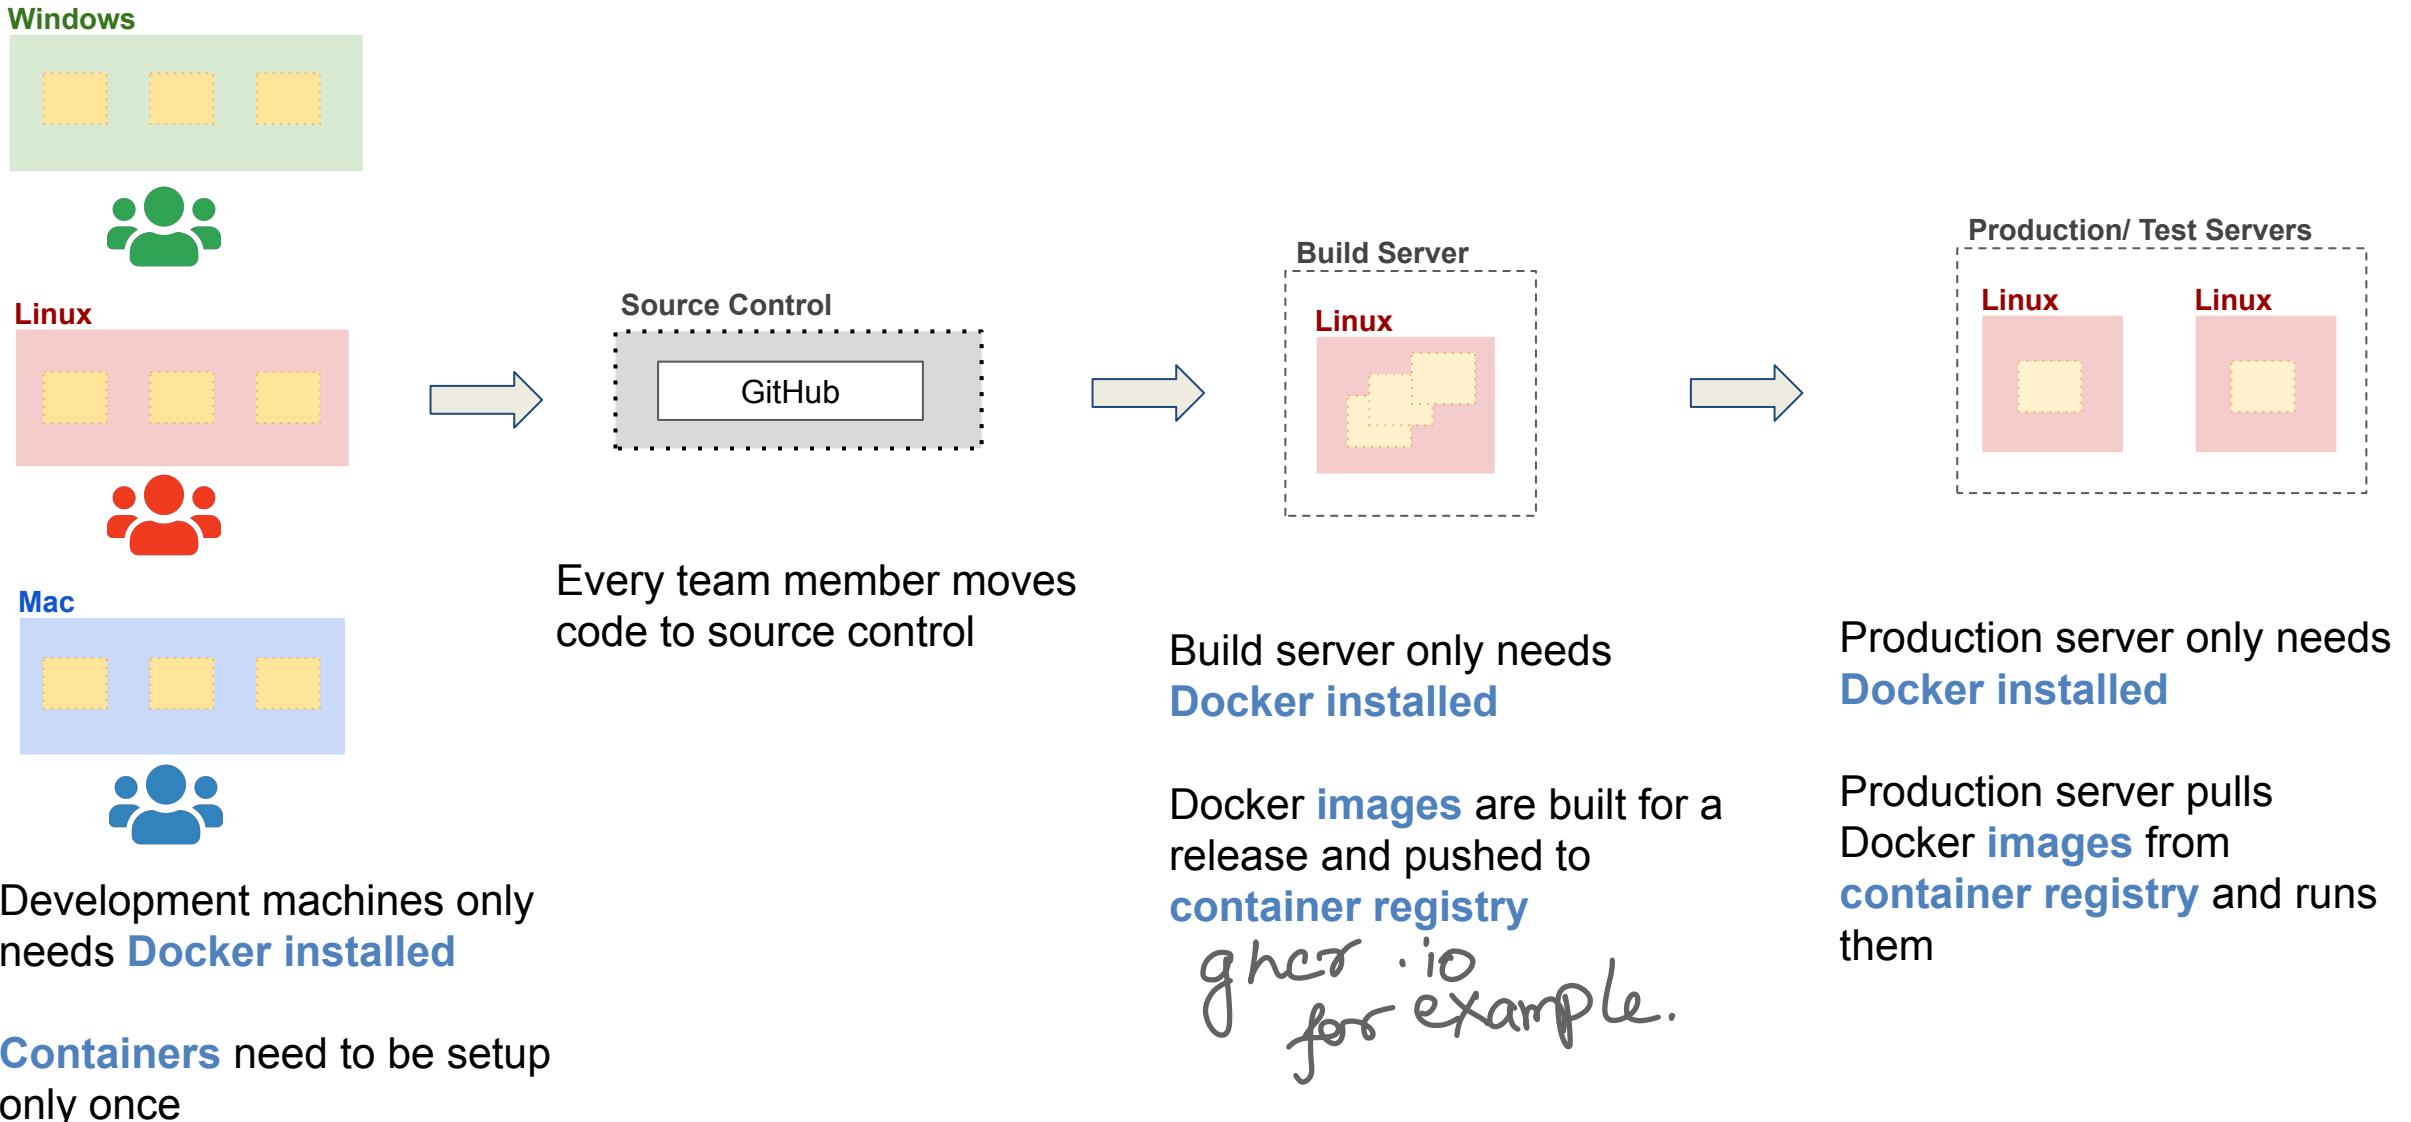
Build server needs to be
installed with all required
softwares/frameworks

Production build is performed
by pulling code from source
control

Production server needs to
be **installed** with all required
softwares/frameworks

Production server will be
different OS version than
development machines

Software Development Workflow (with Docker)



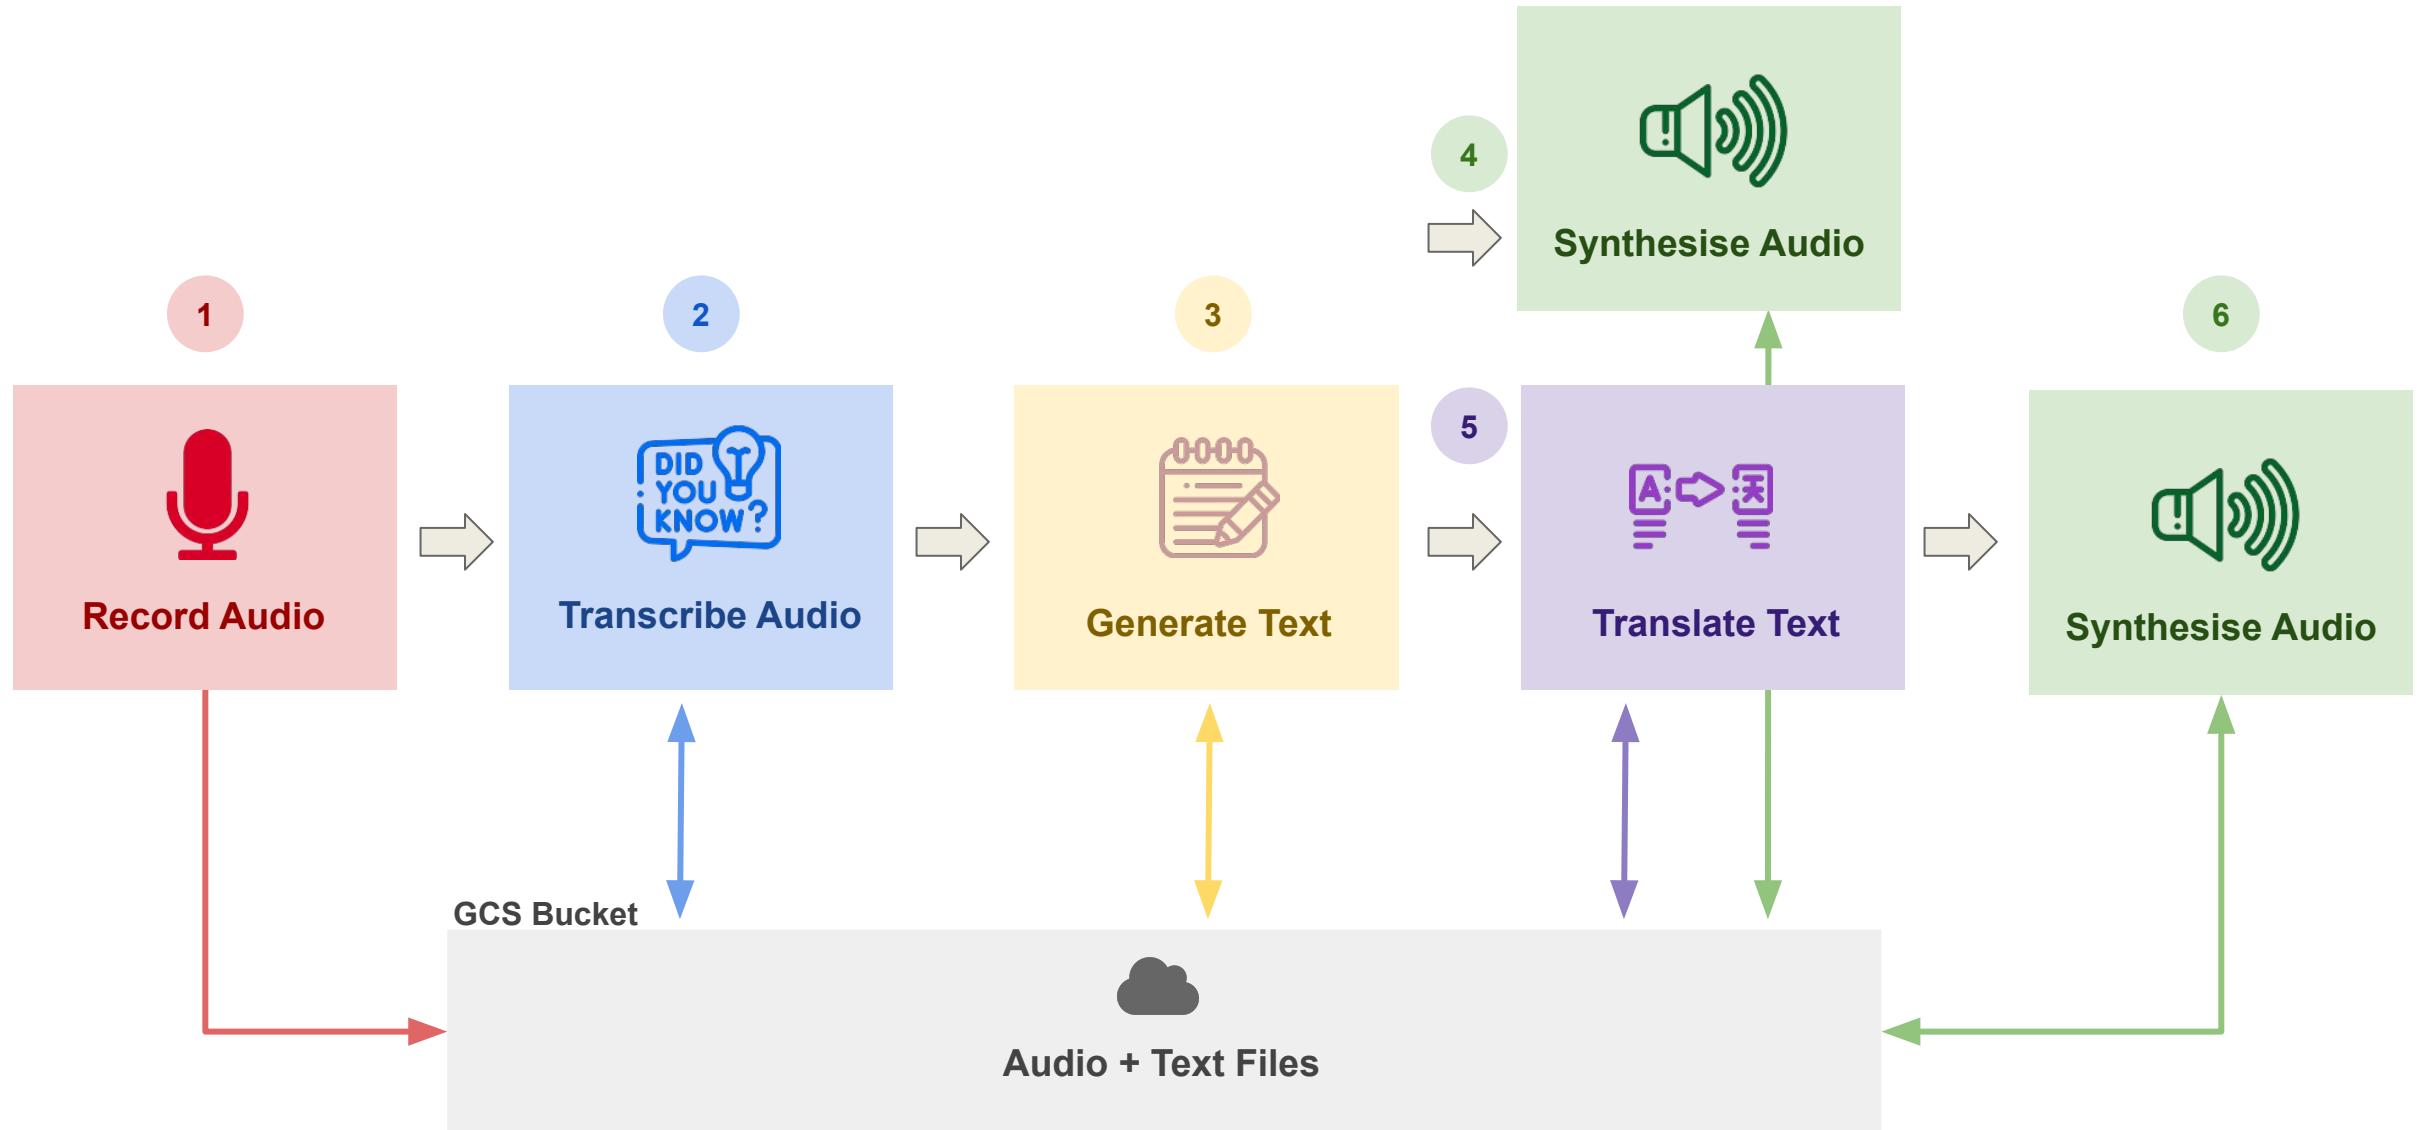
DEVCONTAINERS

- Create docker containers for development : use vscode or jupyter lab.
- Can be enabled on ANY repository on github
- Customize dev experience using "devcontainer.json" standard that builds on docker.
- Use devpod.sh on local machines. Switch between local dev /gpu dev etc based on docker containers locally or on any cloud.

Outline

1. Recap: Review of Previous Material
2. Containers in Architecture: Microservices vs. Monolithic
3. **Implementing Containers as Microservices**

Tutorial - Building the Mega Pipeline App



Tutorial - Building the Mega Pipeline App

≡ AC215: Mega Pipeline App

Click mic to record a Prompt:



Audio Prompts



Transcribed Audio



Generated Text



Synthesised Audio



Translated Text



Synthesised Audio

we will assume that the response variable wide relates to the product of X through some unknown function FX which expresses an underlying

we will assume that the response variable wide relates to the product of X through some unknown function FX which expresses an underlying function, the one with zero sign, that defines what the product-length is, what is the amount of time given to the product, whether the sum is larger or smaller in proportion to the product length, and what is the order in which the product is to be given, and so forth. However, what is a meaningful measure of the product length? It is not always obvious

▶ 0:00 / 0:14 ━━━━ ⏪ ⏹

Nous supposerons que la variable de réponse large concerne le produit de X à travers une fonction de fonction inconnue qui exprime une fonction sous-jacente, celle avec un signe zéro, qui définit la longueur de la longueur du produit, quelle est la quantité de temps donnée au produit, si la somme est plus grande ou plus petite proportionnelle à la longueur du produit et quelle est l'ordre dans lequel le produit doit être donné, etc. Cependant, quelle est une mesure significative de la longueur du produit? Ce n'est pas toujours évident

▶ 0:00 / 0:26 ━━━━ ⏪ ⏹

Tutorial - Building the Mega Pipeline App

- **App:** <https://ai5-mega-pipeline.dlops.io/>
- **Teams**
 -  Team A [transcribe_audio](#):
 -  Team B [generate_text](#):
 -  Team C [synthesis_audio_en](#):
 -  Team D [translate_text](#):
 -  Team E [synthesis_audio](#):
- **Instructions:** <https://github.com/dlops-io/mega-pipeline>

THANK YOU